

Problem 3.9

a) as we can see $\phi(x)$ projected x into \mathbb{R}^n space, for
 $\text{ptA in } \mathbb{R}^n \text{ dim} = n+1$. this is to able to shatter n point.
 which means this transformation works.

b)

$$\phi(x, x_n) = \left(\exp\left(-\frac{\|x - x_n\|^2}{2\sigma^2}\right) \right) \quad (\sigma^2 \text{ is just adding scalar to } x \text{ and } x_n, \text{ thus could ignore it})$$

$$\Rightarrow \exp(-x^2) \exp(-x_n^2) \exp(2xx_n)$$

$$\text{Taylor expansion} = \exp(-x^2) \exp(-x_n^2) \sum_{i=1}^{\infty} \frac{(2xx_n)^i}{i!}$$

$$= \sum_{i=1}^{\infty} \left(\exp(-x^2) \cdot \sqrt{\frac{2^i}{i!}} (x)^i \right) \left(\exp(-x_n^2) \sqrt{\frac{2^i}{i!}} (x_n)^i \right)$$

for known x_n and x .

x is transformed into finite dimensional

with.

$$\left[1, \sqrt{\frac{2}{1!}}, \sqrt{\frac{2}{2!}}, \dots \right] \begin{bmatrix} x^1 \\ x^2 \\ \vdots \\ \vdots \end{bmatrix}$$

for infinity dimension. $\text{dim} = \infty$, and can shatter all x . this
 this transformation works

c) is just a special case of b)

To calculate the euclidean distance, we need to project X to (ij) 's space
When we project X to $\mathbb{R}^{160} \hat{=} \text{proj}(X)$

$\text{proj}(X)$ is just linear combination of elements of X

thus we can use the Taylor expansion around in part b) respect to $\text{proj}(X)$. and transform $\text{proj}(X)$ into infinity dimension. which also means X is transformed to infinity dimension. Thus the transformation works.

Exercise, 4.5

$$a) \quad W'P'PW = (PW)'(PW) = \sum_{q=0}^Q w_q^2$$

$$\text{thus } PW = \begin{bmatrix} w_0 \\ \vdots \\ w_Q \end{bmatrix}_{(Q+1) \times 1} \Rightarrow P = I_{Q+1}$$

$$b) \quad (PW)'(PW) = \left(\sum_{q=0}^Q w_q \right)^2$$

$$PW = \begin{bmatrix} \frac{1}{\sqrt{Q}} \sum_{i=0}^Q w_i \\ \vdots \\ \frac{1}{\sqrt{Q}} \sum_{i=0}^Q w_i \end{bmatrix}_{(Q+1) \times 1} \Rightarrow P = \frac{1}{\sqrt{Q+1}} \begin{bmatrix} 1 & 1 & \dots & 1 \\ \vdots & \vdots & \ddots & \vdots \\ 1 & 1 & \dots & 1 \end{bmatrix}_{(Q+1) \times (Q+1)}$$

problem 4.8

from book Equation (4.6)

$$E_{\text{aug}}(w) = E_{\text{in}}(w) + \lambda w^T w$$

take derivative respect to w on both side.

$$\nabla E_{\text{aug}}(w) = \nabla E_{\text{in}}(w) + 2\lambda w$$

$$\Rightarrow w(t) - \eta \nabla E_{\text{aug}}(w(t)) = w(t) - \eta (\nabla E_{\text{in}}(w(t)) + 2\lambda w(t))$$

$$= w(t) - \eta \nabla E_{\text{in}}(w(t)) - 2\lambda \eta w(t)$$

$$= (1 - 2\lambda \eta) w(t) - \eta \nabla E_{\text{in}}(w(t))$$

CSE 417 homework 3

Muzhou Liu

Problem 1

- The generalization of the model looked okay, as there is small difference between the training error and the test error in each case, which shows the model is well generalized.
 - The E_{in} we got from the 10k, 100k and 1 million iterations are 0.5847, 0.4937 and 0.4354
 - The training error we got from 10k, 100k and 1 million iterations are 0.3092, 0.2237, and 0.1513
 - The test error we got from 10k, 100k and 1 million iterations are 0.3172, 0.2069, and 0.1310
 - As the three interactions all used up maxim iteration steps we set up, the more the algorithm iterates, the closer the gradient to its minimum state. Which is the reason that we can see that in all criteria, the output is improving as the number of iteration increases.
- Comparing with the results we got from last part, glm model, it has the lowest E_{in} (0.4074), little bit higher training error (0.1711) and the lowest testing error (0.1103)
 - Compare the running time using the “timeit” function, it returns 5.1616×10^{-4} second for the ‘glmfit’ function and 7.5506 seconds for the one-million-time iteration which returns an even worse output comparing to the ‘glmfit’ function.
- After standardization, we tried five different learning rates (from 10^{-1} to 10^{-5}) to input into the model, they took 2333, 23368, 233708, 2337116 and 23371191 steps to final terminate.
 - The e_{in} we got from all the five learning are all 0.4074.

Appendix 1 :Codes

```
function [ w, e_in ] = logistic_reg( X, y, w_init, max_its, eta )
%LOGISTIC_REG Learn logistic regression model using gradient descent
% Inputs:
%     X : data matrix (without an initial column of 1s)
%     y : data labels (plus or minus 1)
%     w_init: initial value of the w vector (d+1 dimensional)
%     max_its: maximum number of iterations to run for
%     eta: learning rate

% Outputs:
%     w : weight vector
%     e_in : in-sample error (as defined in LFD)

X = [ones(size(X,1),1) X];
y = 2*y -1;
old_w = w_init;
n_itr = 0;
tol = 0.001*2;
n = size(X,1);
while tol > 0.001 && n_itr < max_its
    g = -sum(y .* X ./ (1+exp(y.*X*old_w)))/n;
    new_w = old_w-eta*transpose(g);
    old_w = new_w;
    n_itr = n_itr+1;
    tol = max(abs(g)) ;
end
w = new_w;
e_in = sum(log(1+exp(-y.*X*w)))/n;
end

function [ test_error ] = find_test_error( w, X, y )
%FIND_TEST_ERROR Find the test_error of a linear separator
% This function takes as inputs the weight vector representing a linear
% separator (w), the test examples in matrix form with each row
% representing an example (X), and the labels for the test data as a
% column vector (y). X does not have a column of 1s as input, so that
% should be added. The labels are assumed to be plus or minus one.
% The function returns the error on the test examples as a fraction. The
% hypothesis is assumed to be of the form (sign ( [1 x(n,:)] * w ))

y = 2*y-1;
temp = ones(size(X,1),1);
X = [temp X];
test_label = sign(X*w);

test_error = sum(y ~= test_label)/size(X,1);
```

```
end
```

```
function [ w, e_in, n_itr ] = logistic_reg2( X, y, w_init, eta )
%LOGISTIC_REG Learn logistic regression model using gradient descent
%   Inputs:
%       X : data matrix (without an initial column of 1s)
%       y : data labels (plus or minus 1)
%       w_init: initial value of the w vector (d+1 dimensional)
%       max_its: maximum number of iterations to run for
%       eta: learning rate

%   Outputs:
%       w : weight vector
%       e_in : in-sample error (as defined in LFD)

X = [ones(size(X,1),1) X];
y = 2*y-1;
old_w = w_init;
n_itr = 0;
tol = 0.001*2;
n = size(X,1);

while tol > 10^(-6)
    g = -sum(y .* X ./ (1+exp(y.*X*old_w)))/n;
    new_w = old_w-eta*transpose(g);
    old_w = new_w;
    n_itr = n_itr+1;
    tol = max(abs(g)) ;
end

n_itr= n_itr;
w = new_w;
e_in = sum(log(1+exp(-y.*X*w)))/n;
end

train = csvread('clevelandtrain.csv',1,0);
train_x = train(:,1:13);
train_y = train(:,14);

test = csvread('clevelandtest.csv',1,0);
test_x = test(:,1:13);
test_y = test(:,14);

[w_10k,e_10k] = logistic_reg(train_x,train_y,zeros(14,1),10000,0.00001);
[w_100k,e_100k] = logistic_reg(train_x,train_y,zeros(14,1),100000,0.00001);
[w_1m,e_1m] = logistic_reg(train_x,train_y,zeros(14,1),1000000,0.00001);
```

```

train_error_10k = find_test_error(w_10k,train_x,train_y)
test_error_10k = find_test_error(w_10k,test_x,test_y)

train_error_100k = find_test_error(w_100k,train_x,train_y)
test_error_100k = find_test_error(w_100k,test_x,test_y)

train_error_lm = find_test_error(w_lm,train_x,train_y)
test_error_lm = find_test_error(w_lm,test_x,test_y)


glm_w = glmfit(train_x,train(:,14) , 'binomial');

e_glm = sum(log(1+exp(-(2*train_y-1).*[ones(size(train_x,1),1)
train_x]*glm_w)))/size(train_x,1)
glm_train_error = find_test_error(glmfit(train_x,train(:,14) ,
'binomial'),train_x,train_y)
glm_test_error = find_test_error(glmfit(train_x,train(:,14) ,
'binomial'),test_x,test_y)

glm_time = @() glmfit(train_x,train(:,14) , 'binomial');
glm_time = timeit(glm_time)
my_time = @() logistic_reg(train_x,train_y,zeros(14,1),1000000,0.00001);
my_time = timeit(my_time)

stand_train_x = zscore(train_x);


[w_std1,e_std1, n_std1] = logistic_reg2(stand_train_x, train_y,
zeros(14,1),0.1)

[w_std2,e_std2, n_std2] = logistic_reg2(stand_train_x, train_y,
zeros(14,1),0.01)

[w_std3,e_std3, n_std3] = logistic_reg2(stand_train_x, train_y,
zeros(14,1),0.001)

[w_std4,e_std4, n_std4] = logistic_reg2(stand_train_x, train_y,
zeros(14,1),0.0001)

[w_std5,e_std5, n_std5] = logistic_reg2(stand_train_x, train_y,
zeros(14,1),0.00001)

%[w_std6,e_std6, n_std6] = logistic_reg2(stand_train_x, train_y,
zeros(14,1),0.000001);

train_error_std1 = find_test_error(w_std1,stand_train_x,train_y)
train_error_std2 = find_test_error(w_std2,stand_train_x,train_y)
train_error_std3 = find_test_error(w_std3,stand_train_x,train_y)

```



```
train_error_std4 = find_test_error(w_std4,stand_train_x,train_y)
train_error_std5 = find_test_error(w_std5,stand_train_x,train_y)
```

Appendix 2 : Outputs

Workspace		
Name ▲	Value	
e_100k	0.4937	
e_10k	0.5847	
e_1m	0.4354	
e_glm	0.6062	
e_std1	0.4074	
e_std2	0.4074	
e_std3	0.4074	
e_std4	0.4074	
e_std5	0.4074	
glm_test_error	0.1103	
glm_time	5.1616e-04	
glm_train_error	0.1103	
glm_w	14x1 double	
my_time	7.5506	
n_std1	2333	
n_std2	23368	
n_std3	233708	
n_std4	2337116	
n_std5	23371191	
stand_train_x	152x13 double	
test	145x14 double	
test_error_100k	0.2069	
test_error_10k	0.3172	
test_error_1m	0.1310	
test_x	145x13 double	
test_y	145x1 double	
train	152x14 double	
train_error_100k	0.2237	
train_error_10k	0.3092	
train_error_1m	0.1513	
train_error_std1	0.1711	
train_error_std2	0.1711	
train_error_std3	0.1711	
train_error_std4	0.1711	
train_error_std5	0.1711	
train_x	152x13 double	
train_y	152x1 double	
w_100k	14x1 double	
w_10k	14x1 double	
w_1m	14x1 double	
w_std1	14x1 double	
w_std2	14x1 double	
w_std3	14x1 double	
w_std4	14x1 double	
w_std5	14x1 double	